

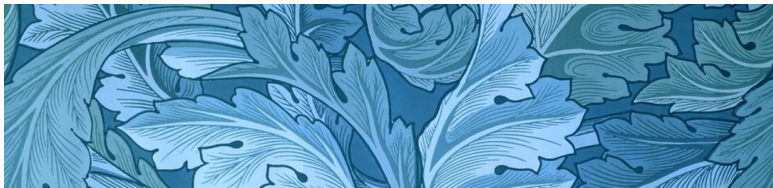


The Holy Grail of Gradual Security

Final Examination for Doctor of Philosophy in Computer Science

Tianyu Chen

Indiana University



Tianyu's Thesis Statement

It *is possible* to design a gradual IFC programming language that satisfies both noninterference and the gradual guarantee while supporting type-based reasoning, by excluding the unknown label ★ from runtime security labels and using security coercions to represent casts.

Road Map

👉 Background

- Explicit flow and implicit flow
 - Information flow control (IFC): static, dynamic, and gradual
 - The gradual guarantee and its tension with IFC
- ▶ Source of the tension: including ★ in runtime labels
 - ▶ Comparing λ_{IFC}^* with GSL_{Ref}
 - ▶ Timeline of dissertation writing

Explicit Information Flow

Can we infer output from input in the following program?

```
let input = private-input () in  
  publish (¬ input)
```

Explicit Information Flow

Can we infer output from input in the following program?

```
let input = private-input () in  
  publish ( $\neg$  input)
```

- ✓ Yes!
 - ▶ Witness at least two executions
 - ▶ Output is the negation of input
 - ▶ Explicit flow

Implicit Information Flow

Can we infer output from input in the following program?

```
let input = private-input () in  
  publish (if input then false else true)
```

Implicit Information Flow

Can we infer output from input in the following program?

```
let input = private-input () in  
  publish (if input then false else true)
```

- ✓ Also yes
 - ▶ Again, output is the negation of input
 - ▶ **Implicit flow**: input influences output through *branching*

Information-Flow Control (IFC)

- ▶ Ensures that information transfers adhere to a security policy
- ▶ For example, **high** input must not flow to **low** output
- ▶ Propagate and check the security labels
- ▶ IFC in PL $\left\{ \begin{array}{l} \text{static} \text{ using a type system} \\ \text{dynamic} \text{ using runtime monitoring} \end{array} \right.$

Static IFC Accepts Legal Explicit Flow

(Static IFC using a type system)

```
1 let fconst = λ b : Boolhigh. false in
2 let input  = private-input () in
3 let result = fconst input in
4   publish result
```

✓ Well-typed and runs successfully to unit

- ▶ Why? The return value of fconst is $\begin{cases} \text{always false} \\ \text{of low-security} \end{cases}$
- ▶ Accepted by type-checker. No runtime check

^oprivate-input : Unit_{low} → Bool_{high} and publish : Bool_{low} → Unit_{low}

Static IFC Rejects **Illegal** Explicit Flow

(Replace fconst with flip)

```
1 let flip    = λ b : Boollow . ¬ b in
2 let input   = private-input () in
3 let result  = flip input in  // compilation error
4   publish result
```

✗ **Ill-typed.** Illegal explicit flow:

- input is **high**
 - flip expects **low** argument
- Rejected by type-checker. Again no runtime check

Dynamic Enforcement of Explicit Flow

(Revisit flip with dynamic IFC)

```
1 let flip      =  $\lambda b. \neg b$  in
2 let input     = private-input () in
3 let result    = flip input in
4   publish result    // runtime error
```

✗ **Errors** at runtime (regardless of input)

- ▶ **A runtime check** happens before calling publish

In dynamic IFC, runtime values are tagged with their security level.
The labels can originate from

- ▶ primitive operations
- ▶ annotations on literals
- ▶ the security level of the execution context

Static Enforcement of Implicit Flow

(Different behavior in different branches)

```
1 let flip : Boolhigh → Boollow =  
2   λ b : Boolhigh. if b then false else true in  
3 let input = private-input () in  
4 let result = flip input in  
5   publish result
```

✗ Ill-typed

- ▶ Security label on the type of if is the join (least upper bound) of its branches (_{low}) and the branch condition (_{high}).
- ▶ Rejected by type-checker. No runtime check

Dynamic Enforcement of Implicit Flow

(Enforcing implicit flow with dynamic IFC)

```
1 let flip    = λ b. if b then false else true in
2 let input  = private-input () in
3 let result = flip input in
4   publish result
```

✗ **Errors** at runtime (regardless of input)

- ▶ flip produces a **high** value because of **high** branch condition
- ▶ A runtime check happens before calling publish
- ▶ Illegal implicit flow ruled out **at runtime**

Gradual Typing Bridges Static and Dynamic IFC

Partially-annotated flip:

```
1 let flip : Bool★ → Boollow =  
2   λ b : Bool★. if b then false else true in  
3 let input = private-input () in  
4 let result = flip input in  
5   publish result
```

- ▶ Well-typed but errors at runtime
- ▶ Checking happens on the boundaries between static and dynamic fragments
- ▶ The information flow violation is detected earlier than the dynamic version, as flip returns

The Gradual Guarantee

less precise

```
let f : Bool★ → Bool★ =  
  λ b : Bool★. true in  
let i = private-input () in  
let result = f i in  
publish result
```

```
let f : Bool★ → Boollow =  
  λ b : Bool★. true in  
let i = private-input () in  
let result = f i in  
publish result
```

more precise

```
let f : Boolhigh → Boollow =  
  λ b : Boolhigh. true in  
let i = private-input () in  
let result = f i in  
publish result
```

- ▶ In the absense of errors, adding or removing security annotations does not change the result of the program
- ▶ Adding security annotations may trigger errors
- ▶ Removing security annotations may not trigger errors

Static Enforcement of Flows Through Mutable References

```
1 let a      = ref low true in
2 let input = private_input () in
3 if input then
4     a := false
5 else
6     a := true
7 publish (! a)
```

- ▶ The reference has type `Ref (Boollow)`. It points to a low memory location
- ▶ The type of the branch condition is `Boolhigh`
- ✗ ~~Writing to low memory under a high branch condition~~

Dynamic Enforcement of Flows Through Mutable References

```
1 let a      = ref low true in
2 let input = private_input () in
3 if input then
4     a := false
5 else
6     a := true
7 publish (! a)
```

The assignments fail at runtime because the no-sensitive-upgrade (NSU) mechanism¹ prevents writing to a **low** security pointer in a **high** security branch.

¹ Austin and Flanagan. *Efficient purely-dynamic information flow analysis*. PLAS 2009.

Counterexample of Gradual Guarantee in GSL_{Ref}

less precise

```
1 let x = private-input () in
2 let a = ref ★ true★ in
3 if x then (a := falsehigh)
4       else ()
```

more precise

```
let x = private-input () in
let a = ref hightruehigh in
if x then (a := falsehigh)
       else ()
```

- ✓ The **more precise** program (right) runs successfully
- ✗ But the **less precise** version (left) errors in GSL_{Ref} ²
 - ▶ The assignment fails because it is in a high-security branch and GSL_{Ref} conservatively treats the reference's label (★) as if it were **low**

²Toro, Garcia, Tanter. *Type-Driven Gradual Security with References*. TOPLAS 2018.

But wait... GSL_{Ref} allows \star labels on values?

The counterexample depends on labeling a reference with unknown security (\star):

```
1 let x = private-input () in
2 let a = ref  $\star$  true $\star$  in
3 if x then (a := falsehigh)
4         else ()
```

- ▶ Dynamic IFC languages don't use \star as a runtime security label.
- ▶ Gradual languages traditionally use \star for type checking, but not for categorizing runtime values.
- ▶ The inputs to an information flow system are the user's choices regarding what data is high or low security.

Sources of the Tension with the Gradual Guarantee

Lang.	Noninterference	Gradual Guarantee	Type-guided classification	NSU	Runtime security labels
GSL _{Ref}	✓	✗	✓	✓	{low, high, ★}
GLIO	✓	✓	✗	✓	{low, high}
WHILE ^G	✓	✓	✓	✗	{low, high, ★}
λ_{Ifc}^* (ours)	✓	✓	✓	✓	{low, high}

Removing ★ from the runtime labels enables the gradual guarantee.

λ_{IFC}^* Excludes ★ From Runtime Labels

less precise

```
1 let x = private-input () in
2 let a : (Ref Bool★)★ =
3   ref high truehigh in
4 if x then (a := falsehigh)
5   else ()
```

more precise

```
let x = private-input () in
let a : (Ref Boolhigh)high =
  ref high truehigh in
if x then (a := falsehigh)
  else ()
```

- ✓ The **more precise** program runs **successfully** to unit
- ✓ The **less precise** program also runs **successfully** to unit

λ_{IFC}^* Excludes \star From Runtime Labels

less precise

```
1 let x = private-input () in
2 let a : (Ref Bool $\star$ ) $\star$  =
3   ref high truehigh in
4 if x then (a := falsehigh)
5   else ()
```

more precise

```
let x = private-input () in
let a : (Ref Boolhigh)high =
  ref high truehigh in
if x then (a := falsehigh)
  else ()
```

- ✓ The more precise program runs successfully to unit
- ✓ The less precise program also runs successfully to unit
- ✓ Problem solved!



Comparing λ_{IFC}^* With GSL_{Ref}

- ▶ The default security label in λ_{IFC}^* is **low**, so the programmer does not have to label constants
- ▶ Remove the labels on constants to obtain the following program, which also reduces successfully to unit:

```
let x = private-input () in
let a : (Ref Bool★)★ = ref high true in
if x then (a := false)
  else ()
```

Comparing with the program in GSL_{Ref} , which errors:

```
let x = private-input () in
let a = ref ★ true★ in
if x then (a := falsehigh)
  else ()
```

The Syntax of $\lambda_{\text{IFC}}^{\star}$

We define a gradual IFC calculus $\lambda_{\text{IFC}}^{\star}$ with mutable references, first-class functions, and conditionals.

Highlighted security labels default to **low** if omitted:

$$\begin{aligned}\ell &\in \{\text{low}, \text{high}\} \\ g &\in \{\text{low}, \text{high}, \star\} \\ \iota &::= \text{Unit} \mid \text{Bool} \\ T &::= \iota \mid A \xrightarrow{g} A \mid \text{Ref } (T_g) \\ A &::= T_g \\ M &::= x \mid (\$ k)_{\ell} \mid (\lambda^g x:A. M)_{\ell} \mid (M M)^p \\ &\quad \mid (\text{if } M \text{ then } M \text{ else } M)^p \\ &\quad \mid (\text{ref } \ell M)^p \mid !^p M \mid (M := M)^p\end{aligned}$$

Vigilance: Type-Based Reasoning for Explicit Flows

Consider the example from Toro et al. [2018]:

```
1 let mix : Intlow → Inthigh → Intlow =  
2   λ pub priv .  
3     if pub < (priv : Int★ : Intlow) then 1 else 2 in  
4 mix 1low 5low
```

Free theorem: The mix function either ① returns a value that does not depend on priv or ② produces a runtime error

In λ_{IFC}^* , $5 \langle \uparrow; \text{high}!; \text{low}?^p \rangle \longrightarrow \text{blame } p$

Type-Guided Classification:

Type-Based Reasoning for Implicit Flows

Consider another example from Toro et al. [2018]:

```
1 let mix : Intlow → Int* → Intlow =  
2   λ pub priv. if pub < priv then 1 else 2 in  
3 let smix : Intlow → Inthigh → Intlow =  
4   λ pub priv. mix pub priv in  
5 smix 1low 5low
```

Free theorem: The `smix` function either ① returns a value that does not depend on `priv` or ② produces a runtime error

Type-Based Reasoning for Implicit Flows in λ_{IFC}^*

```
let mix : Intlow → Int* → Intlow =  
  λ pub priv. if pub < priv then 1 else 2 in  
let smix : Intlow → Inthigh → Intlow =  
  λ pub priv. mix pub priv in  
smix 1low 5low
```

```
let mix = λ pub priv.  
  (if ((pub < low!) < priv)  
    then (1 < low!)  
    else (2 < low!))) < low?p > in  
⇒ let smix = λ pub priv.  
  mix pub (priv < high! >) in  
  smix 1 (5 < ↑ >)
```

→* (if (1 < low! > < 5 < ↑; high! > >) then 1 < low! > else ...) < low?^p >

→* (if (true < ↑; high! > >) then 1 < low! > else ...) < low?^p >

→* (prot high (1 < low! >)) < low?^p >

→* 1 < ↑; high! > < low?^p >

→* blame *p*

The Cast Calculus λ_{IFC}^c

The semantics of $\lambda_{\text{IFC}}^\star$ is by translation \mathcal{C} to a cast calculus λ_{IFC}^c .

The casts are represented by coercions on types (a la Henglein) and coercions on security labels.

$$\begin{aligned}
 c &::= \text{id}(g) \mid \uparrow \mid \ell! \mid \ell?^p \mid \perp^p \\
 \bar{c} &::= \text{id}(g) \mid \perp^p g_1 g_2 \mid \bar{c}; c \\
 e &::= \ell \mid \text{blame } p \mid e \langle \bar{c} \rangle \\
 c_r &::= \text{id}(\iota) \mid \text{Ref } c \ c \mid (\bar{c}, c \rightarrow c) \\
 \mathbf{c} &::= c_r, \bar{c} \\
 M &::= x \mid \$k \mid \text{addr } n \mid \lambda x. M \mid \text{let } x=M:A \text{ in } M \\
 &\mid \text{app } M \ M \ A \ B \ \ell \mid \text{app}^\star M \ M \ A \ T \\
 &\mid \text{if } M \ A \ \ell \ M \ M \mid \text{if}^\star M \ T \ M \ M \\
 &\mid \text{ref } \ell \ M \mid \text{ref}?^p \ell \ M \mid ! M \ A \ \ell \mid !^\star M \ T \\
 &\mid \text{assign } M \ M \ T \ \ell \ \ell \mid \text{assign}?^p M \ M \ T \ g \\
 &\mid \text{prot } e \ \ell \ M \ A \mid M \langle c \rangle \mid \text{blame } p
 \end{aligned}$$

Conclusion

- ▶ Noninterference and the gradual guarantee can co-exist if we keep \star out of the runtime security labels.
- ▶ The security labels on constants and memory locations should default to **low** or **high** instead of \star .
- ▶ The Agda mechanization of $\lambda_{\text{IFC}}^{\star}$ is at <https://github.com/Gradual-Typing/LambdaIFCStar>

Thank you! 😊